

The Hidden Variable in Embedded VLM Perception: Temporal Calibration for Real-Time Work-Zone State Estimation

Anonymous Authors
Paper under double-blind review

Abstract—Vision-language models (VLMs) are increasingly proposed as replacements for detector pipelines in driving perception, usually compared against the baseline’s published temporal smoothing constants. We show on embedded hardware that this hidden variable can decide the outcome. We run an INT4 2B VLM fine-tuned on $\sim 46\text{K}$ work-zone question-answer pairs and $\sim 10\text{K}$ hard negatives and a zero-shot 4B physical-AI reasoner on a Jetson AGX Orin through an evidence-first prompt cascade, and compare them with a YOLO+CLIP pipeline under a latency-honest protocol in which measured latency drives frame dropping. On 484K annotated frames of ROADWork driving, constants inherited from the 2B leave the 4B model at macro-F1 0.290; recalibrated, it reaches 0.515. Recalibration lifts the detector baseline from 0.470 to 0.551, the best result, and no VLM-based system then beats it; a counted hidden-Markov filter on the detector score alone matches it without a threshold search. These conclusions hold on drive-disjoint folds, at matched false-alarm rates and with matched search budgets, and live concurrent execution on the Orin reproduces them.

I. INTRODUCTION

Work zones concentrate difficult driving-perception conditions: their appearance is regional and unstandardized, geometry changes daily, evidence such as cones, barrels, signs, and workers is small or occluded, and both misses and nuisance alerts are costly. They accounted for roughly 101,000 U.S. crashes and 899 fatalities in 2023 [1], yet production ADAS rarely model work zones explicitly and temporary speed limits are often absent from digital maps.

The strongest onboard approach combines a YOLO-style work-zone detector, CLIP verification [2], exponential smoothing, and hysteresis to estimate four traversal states [3], [4]. It is fast and effective after a zone is established, but reported IoU falls from 0.56/0.47 on OUTSIDE/INSIDE to 0.11/0.09 on APPROACHING/EXITING [3]. These boundaries are precisely where an advisory is most useful: an alert after the vehicle reaches the barrels has little anticipatory value.

Meanwhile, fine-tuned VLMs describe work-zone scenes well offline [5], and weight-only quantization plus embedded runtimes bring small VLMs to automotive hardware [6], [7]. Whether they should replace detector pipelines is usually judged by comparing systems as published, each with the smoothing and hysteresis constants its authors tuned. We build the strongest embedded VLM estimator we can, run it beside the detector pipeline on the same device, videos and protocol, and ask what such a comparison actually measures.




Fig. 1. The hidden variable. Validation macro-F1 (208 videos) under three comparison regimes. Left: every VLM runs the 2B’s temporal constants and the detector pipeline its published ones; the zero-shot 4B reasoner (C3E) looks useless. Middle: the VLM systems are calibrated and the baseline is not, as in a typical comparison; C3E and the joint estimator beat the detector pipeline. Right: the baseline is calibrated by the same procedure and becomes the best system.

The answer is: mostly the temporal constants (Fig. 1). Against the hand-tuned detector pipeline our VLM-based systems look competitive or better, but constants inherited from one checkpoint leave a 4B reasoner at INSIDE IoU 0.007, and its own calibration lifts it to 0.574. The same calibration raises the detector pipeline from 0.470 to 0.551, after which no VLM-based system beats it. A team choosing a perception stack from the usual comparison would adopt a costlier VLM stack for no measurable gain.

Our contributions are three. **(1) Evaluation:** a latency-honest protocol, confirmed by live concurrent execution, with per-model recalibration by evidence replay. **(2) Findings:** temporal constants do not transfer between the systems we test, and calibrating the baseline reverses the ranking; this holds on drive-disjoint folds, at matched false alarms and with matched search budgets. **(3) Systems:** a real-time evidence-first cascade for a quantized 2B VLM, an INT4 Cosmos3-Edge reasoner on the same Orin, and a counted hidden-Markov filter that matches the recalibrated baseline from the detector score alone, with an analysis that locates the residual VLM gap in ego-relevance.

II. RELATED WORK

a) Work-zone perception.: Early systems used environmental models with sensor fusion [8] or frame classification [9]. ROADWork [5] provides 4,375 videos and 9.6K an-

notated images from 18 U.S. cities and shows large gains from in-domain fine-tuning (+32.2 AP, +36.7 SPICE). The closest system [3] deploys a ROADWork-trained YOLO detector with CLIP verification, EMA smoothing and hysteresis, and defines the traversal-state evaluation we adopt.

b) *VLMs and world models for driving.*: VLMs support planning [10], [11], description [12] and reasoning [13], and Cosmos world models plan at datacenter or automotive-SoC scale [14], [15]; we estimate one safety-relevant state under an embedded latency budget.

c) *Efficient inference and cascades.*: AWQ [6] and embedded runtimes [7] enable multi-Hz VLM inference on embedded GPUs. Detection [16] and LLM [17] cascades chain models of increasing cost and escalate on confidence; our stages share one model, and evidence semantics (transcribed text versus a binary answer) drive escalation.

d) *Temporal filtering and fair comparison.*: Our joint estimator is a discrete hidden Markov model [18] run as a Bayes forward filter [19]; we claim no novelty for it. That baselines deserve the proposed method’s tuning budget is known [20]; we measure how much it changes conclusions for embedded temporal estimators.

III. SYSTEM

A. Two embedded deployments

We deploy *two* models on the same Jetson AGX Orin 64 GB, which shares its SoC with the automotive DRIVE AGX Orin, adapted in the two ways such a model admits: one by task fine-tuning, one by quantization and per-model calibration alone. Both run the identical prompt channels, cascade, protocol and splits, so every comparison in this paper varies the model and its adaptation, never the pipeline around it.

Fine-tuned 2B. A 2B-parameter VLM (Qwen3-VL architecture [21], initialized from Cosmos-Reason2-2B [22]), used for perception only, fine-tuned through a curriculum derived from ROADWork [5], whose imagery spans 18 U.S. cities under one signing standard (MUTCD [23]): work-zone VQA (~46K pairs), ego-state supervision, sign reading, a general-driving retention mix, and ~10K work-zone-free BDD100K hard negatives. The last stage brought the existence gate from 95% to 0% “yes” on clean BDD100K scenes (on ROADWork OUTSIDE seconds, which include visible but irrelevant work, it is less selective, Table IV).

The language backbone is quantized to INT4 (W4A16 AWQ, group size 128) and the vision tower kept in FP16, both as on-device TensorRT engines [7] (1.38 + 0.83 GB). Decoding (10.4 ms/token) dominates cost, followed by vision encoding (44.6 ms at 720p, ~20 ms at our 480px).

Quantized 4B world model. Cosmos3-Edge (C3E) [14] pairs an autoregressive reasoner with a trajectory-generating diffusion transformer; we run only the reasoner, and “world model” below refers to it. It is quantized to W4A16 INT4 (1.27 GB) with an FP16 visual tower (983 MB) on the same Orin; its frozen 26×46 patch grid fixes inputs at 736×416, and channel latencies are 162/182/488/556 ms (GATE/EGO/DESC/SIGN). It has seen no work-zone training




Fig. 2. Evidence-first cascade. In OUTSIDE, GATE and SIGN run every cycle and candidate evidence invokes DESC. Entry requires a sustained, corroborated vote or high-confidence same-cycle evidence. In active states, EGO drives a Bayesian filter; sustained gate clearance returns to OUTSIDE, while partial fading triggers EXITING. Constants are calibrated per model: 5/3/4/2 for the 2B and 3/3/2/1 for C3E, so the structure transfers but the numbers do not. Diagram values are the 2B’s, with rounded latencies (measured values in Section V-D).

data. Its engines required six patches to the embedded runtime, released with the code.

B. Prompt channels and their failure modes

We operate four fixed-prompt channels with fixed token budgets. Because the first generated token is an answer-start token, the minimum useful budget is two tokens.

a) *Affirmation bias.*: The GATE asks a neutral existence question, but binary prompts cannot be specialized by rephrasing: path-relevance variants all answered “yes” on OUTSIDE probes and distance prompts “close” regardless of scene. Binary outputs are therefore weak evidence requiring a temporal vote, never sole triggers. At ~10 ms per output token, a two-token answer costs 75 ms and a description ~500 ms, so evidence quality, not only confidence, gates smoothing.

b) *Faithful transcription.*: The SIGN channel returns verbatim text such as “SHOULDER WORK” and correctly reports absence, while a binary rephrasing misses signs that transcription catches; we match transcripts against a work-zone lexicon.

c) *Structured corroboration.*: DESC often verbalizes disqualifying context (“barriers on right sidewalk”, “work vehicle parked”), so corroboration requires a work-zone object without such a qualifier. Prompts, budgets and lexicons are released with the code.

C. Evidence-first cascade

The cascade (Fig. 2) allocates compute by state.

Outside (patrol). GATE and SIGN run every cycle (2B at 480px: p95 259 ms, 3.9 Hz), and either can invoke DESC. APPROACHING is entered by vote (gate active in K_{enter} of the last N cycles plus one corroboration) or by fast entry on a lexicon-matching sign or same-cycle gate+DESC agreement, which avoids a delay of up to N cycles (~1.3 s for the 2B) for evidence such as “ROAD WORK AHEAD.”

Active states. EGO updates a Bayesian filter whose argmax is restricted to the active states, so that one “outside” reply right after a sign-triggered entry cannot erase it. Only sustained gate clearance returns to OUTSIDE; partial fading triggers EXITING, which EGO never produced. DESC refreshes every 1.5 s. The 2B’s thresholds ($N=5$, $K_{\text{enter}}=3$, $K_{\text{exit}}=4$, $K_{\text{fade}}=2$) were

fixed on calibration data; C3E gets its own (Section VI). With EGO disabled, as selected for C3E, APPROACHING becomes INSIDE on the transition prior alone, 13 cycles after entry.

D. Joint estimator over both evidence streams

Grafting the detector and the world model onto each other’s temporal machinery does not work: the ported detector counts cycles (150-cycle maximum approach, 20 minimum frames outside) tuned for its native loop, and the cascade votes over a window tuned for the 2B’s ~ 260 ms cycle, while C3E answers only every ~ 1.5 s. Either spine imports constants that are wrong for the combination.

The joint estimator therefore inherits no constant: it is a discrete hidden Markov model over the four states, run as a forward filter [18] with maximum-likelihood tables and no EMA, hysteresis threshold or cycle counter. Its observation combines the detector’s frame score (six buckets), the world model’s gate, sign-keyword and corroboration answers ordered by strength (six levels), and the age of that answer (three buckets): 108 cells. The emission is counted once per detector cycle on the calibration split with Laplace smoothing, jointly rather than factorized, because the interaction is the signal: the world model’s APPROACHING is reliable when the detector corroborates it and mostly a false alarm when it does not (73% of those frames are OUTSIDE). The transition is a per-second rate matrix, counted as ground-truth transitions over dwell time and scaled by each step’s dt , so it does not depend on the loop rate. The log-likelihood is scaled by dt/τ ($\tau = 0.5$ s): a world-model answer persists across ~ 30 detector frames, and counting it once per frame drove an early version to 135 false alarms per hour. Bucket edges and τ were set on calibration data (Section VII).

The detector runs at its native rate and is never blocked; the world model is re-pollled as soon as it answers, and its last answer is held with its age. We record raw evidence with measured clocks once per video and replay the filter on it, so ages and frame drops in every result are measured ones; Section VII confirms the replay against live concurrent execution.

IV. EVALUATION PROTOCOL

a) *Benchmark.*: We use the detector baseline’s 520-video ROADWork traversal benchmark [3]: 4.37 h of 1920×1080 , 30 fps driving from 25 drives, with 484K interval-labeled frames. A state-stratified 60/40 split assigns all design choices to calibration (312 videos, 2.59 h) and results to validation (208 videos, 192K frames, 1.78 h, 13 pure-negative). All 25 drives occur in both partitions, so Section VIII repeats every replay-fitted comparison with drive-disjoint folds. ROADWork samples its annotated images from the same drives as its videos, so the 2B’s fine-tuning images, like the baseline detector’s training images, may come from benchmark drives; C3E is zero-shot. Every calibration effect we report compares a system with itself, perception fixed, and is not inflated by this overlap.

TABLE I

VALIDATION, 208 VIDEOS, ONE PROTOCOL AND DEVICE (%; TIMING IN S). DET.+CLIP: HAND-TUNED [3] AND RECALIBRATED (SECTION VI); DET. *cnt.*: DETECTOR SCORE IN THE COUNTED FILTER; 2B: GREEDY, NO POST-FILTER (RECALIBRATED: 0.459); JOINT: DETECTOR+C3E IN THE COUNTED FILTER. LEAD: INSIDE ONSETS PRECEDED BY ≥ 2 S OF ACTIVE STATE. F1@20: AT ≤ 20 FALSE ALARMS/H (SECTION VIII).

Metric	Det.+CLIP		Det.	2B	C3E	Joint
	hand	cal.	cnt.			
Frame acc.	62.9%	64.8%	68.2%	62.0%	65.5%	66.1%
Macro F1	0.470	0.551	0.524	0.453	0.515	0.546
IoU OUT	0.617	0.627	0.652	0.581	0.615	0.641
IoU APPR	0.124	0.296	0.219	0.152	0.278	0.246
IoU IN	0.544	0.552	0.573	0.532	0.574	0.578
IoU EXIT	0.106	0.153	0.124	0.062	0.071	0.160
IN ev. recall	95.7	77.8	88.6	97.3	87.0	85.9
IN ev. prec.	79.9	92.5	75.9	56.8	71.5	66.8
False alarms/h	30.9	20.2	22.5	56.2	64.6	25.9
Entry MAE	2.52	2.31	2.44	2.52	2.15	2.22
IN entry med.	3.47	2.10	2.73	2.67	2.50	2.27
Lead ≥ 2 s	79.4	78.8	73.1	76.2	86.2	66.9
F1@20	0.410	0.529	0.520	0.409	0.487	0.537

b) *Latency-honest camera simulation.*: After each measured inference cycle of duration Δ on the target hardware, the replay advances by Δ and drops intervening frames as a live camera would, so slower systems see sparser frames and pay latency inside the accuracy result. Predictions are held between cycles and scored per frame.

c) *Metrics.*: We adopt the baseline’s suite [3]: frame accuracy and macro-F1; per-state IoU; event precision/recall (an interval is correct if it overlaps one of the same state); entry offsets (Entry MAE over all states, IN entry med. for INSIDE); and false alarms per hour (predicted active intervals overlapping no annotated one), plus a lead metric for timing (Section V-A). Intervals are paired 95% cluster-bootstrap intervals (10,000 replications) that resample drives.

V. THE VLM CASCADE AGAINST THE PUBLISHED BASELINE

A. Comparison with the deployed detector system

Table I places every system on the same videos, device and protocol; this section compares the 2B with the detector pipeline as published. With greedy decoding, our default (Section V-C), the 2B is statistically tied with it on frame accuracy and macro-F1 (-0.018 , $[-0.041, +0.003]$) and matches its INSIDE event recall, but trails on event precision (56.8% vs. 79.9%) and false alarms (56.2 vs. 30.9 per hour): a competitive detector-free result, not a replacement. A causal 1 s INSIDE debounce raises precision to 76.0% (macro-F1 0.478) without retraining. Entry MAE is equal (2.52 s), and the 2B’s INSIDE entry median is 0.80 s better.

Both pipelines handle established INSIDE intervals well and EXITING poorly. The 2B’s APPROACHING IoU is higher (0.152 vs. 0.124), as text and distant evidence appear before detector counts stabilize, a gap the recalibrated detector closes (Section VI).

TABLE II

SINGLE-TOGGLE ABLATIONS OF THE 2B CASCADE (VALIDATION, 208 VIDEOS). LEAD: INSIDE ONSETS PRECEDED BY ≥ 2 S OF ACTIVE STATE.

Variant	Acc.	F1	IN prec.	FA/h	Lead
Cascade, sampled	62.5	0.464	54.1	60.1	79.4
greedy decoding	62.0	0.453	56.8	56.2	76.2
binary SIGN	52.5	0.407	46.1	305.8	86.2
no fast entry	63.1	0.462	54.3	34.3	69.4
no qualifier filter	60.3	0.447	53.9	69.7	81.2
no SIGN channel	63.2	0.456	53.4	47.8	82.5
unrestricted filter	61.4	0.445	53.1	69.1	73.8
base checkpoint	48.8	0.289	—	29.2	—

B. Operational false-alarm characterization

Over 1.78 h the 2B raises 56.2 false active-state events per hour, against 30.9 for the detector pipeline; most are the ego-relevance errors of Section IX, not clean-scene noise. Of 156 matched INSIDE entries, 80.8% fire early (median 2.9 s), which an advisory prefers. Section VIII compares all systems at matched false alarms.

C. Ablations and design validation

Table II toggles one component at a time; five of seven change macro-F1 by at most 0.02. The sign channel must transcribe: an equal-cost binary question cuts accuracy from 62.5% to 52.5% and raises false alarms fivefold. Greedy decoding costs 0.5 accuracy points but improves precision and alarms, so we use it as the reproducible default. The qualifier filter adds 2.2 points and removes 9.6 alarms per hour. Fast entry buys anticipation, 10 more points of onsets warned ≥ 2 s ahead, for 26 more alarms per hour. The base checkpoint collapses under the 2B’s thresholds (Section VI).

D. Operational cost

At 480px, p50/p95 latency is 83/90 ms (GATE), 136/149 (EGO), 155/169 (SIGN) and 468/508 ms (DESC); patrol and active cycles run at p95 259 ms (3.9 Hz) and 239 ms (4.2 Hz), the heaviest path at 768 ms. “Real time” here means keeping pace with a live camera (a 259 ms cycle is 7.8 m at 30 m/s), not meeting an alert deadline. Above idle, the VLM draws 28.6 W and 3–26 J per cycle, the detector 22.4 W and 0.8 J at 24 ms.

VI. CALIBRATING EVERY SYSTEM

Section V ran each VLM with thresholds fitted to the 2B and the detector pipeline with its published constants. Here every system is recalibrated by the same replay procedure, and the ranking reverses (Fig. 1).

a) *Temporal hyper-parameters do not transfer.*: With the 2B’s constants, the zero-shot C3E almost never reaches INSIDE (IoU 0.007, recall 3.8%) and looks useless. Recalibrating the same four integers lifts it to macro-F1 0.515 and INSIDE IoU 0.574, with model, prompts and videos unchanged (Fig. 3). The search covers $N \in \{3, 5, 7, 9\}$, $K_* \in 1..N$ and the EGO/SIGN switches (4,896 configurations, later extended to $N \in \{1, 2, 4\}$ without changing the optimum), selected on




Fig. 3. Ground truth and five systems on one validation video, chosen to illustrate the effect; thumbnails mark the middle of each annotated state. With the 2B’s constants, C3E never leaves APPROACHING; recalibrated, it tracks the zone. The hand-tuned detector pipeline enters INSIDE 5 s early and cuts EXITING short; recalibrated, it follows all four states.

calibration macro-F1 over the 100 videos with recorded VLM evidence. The optimum ($N=3$, $K=3/2/1$) is shorter and less hysteretic and disables EGO (Section IX). Because each channel is recorded once and the state machine replayed, a sweep costs CPU-seconds instead of GPU-hours.

The same search moves the fine-tuned 2B by only +0.006 (0.453→0.459): it rescues a model whose constants were wrong rather than rewarding flexibility.

b) *The baseline too.*: The detector pipeline deserves the same treatment. We record every component on every cycle of all 520 videos (detector counts before and after CLIP verification, global CLIP and colour-context scores) and replay the complete pipeline; its hand-tuned configuration reproduces the live run (0.473 vs. 0.470). Searching its thresholds, EMA rate and duration limits on the 312 calibration videos (three grids of 4,800 configurations) raises validation macro-F1 from 0.470 to 0.551 (+0.077, [+0.039, +0.108]), with 92.5% IN-SIDE precision at 20.2 false alarms per hour: the strongest system we measure. On the detector score alone, without CLIP, the same search gives 0.435→0.530.

The procedure also isolates what fine-tuning buys. Recalibrated on its own evidence, the base checkpoint, identical except for the work-zone curriculum, reaches 0.383 on the search videos against 0.496 for the fine-tuned model: fine-tuning is worth +0.113 macro-F1, not the collapse (INSIDE IoU 0.012) the base model shows under the fine-tuned model’s thresholds.

Every choice is made on the 312 calibration videos and none on validation: prompts and lexicons on the full partition, VLM thresholds on the 100-video subset with recorded VLM evidence, detector thresholds and counted tables on all 312. The budgets therefore favour the detector; given the VLM’s budget, one grid on the same 100 videos, the recalibrated pipeline still reaches 0.522, +0.052 over its hand-tuned version ([+0.038, +0.070]) and level with C3E.

c) *No configuration dominates.*: Calibrated, the world model and the detector pipeline tie on macro-F1 (−0.036,

$[-0.080, +0.005]$) and fail differently: C3E leads INSIDE IoU and entry MAE, the detector pipeline precision and nuisance (20.2 vs. 64.6 false alarms/h). Section VII combines them.

VII. COMBINING THE DETECTOR AND THE WORLD MODEL

The joint estimator reaches 0.546 at 25.9 false alarms per hour. It beats the hand-tuned pipeline ($+0.076, [+0.038, +0.104]$) and the calibrated world model ($+0.031, [+0.006, +0.058]$, with 38.8 fewer alarms per hour), and ties the recalibrated pipeline ($-0.005, [-0.046, +0.033]$), trading INSIDE precision (66.8% vs. 92.5%) for recall (85.9% vs. 77.8%).

Its tables are counted on all 312 calibration videos (5-fold cross-validation: 0.535, s.d. 0.019) and its few design choices behave predictably: the evidence rate τ trades accuracy for nuisance (0.538 at 44.7/h for $\tau=0.25$, 0.504 at 22.0/h for $\tau=1$; we keep 0.5), 20 of the 108 cells carry 97% of the observations, and 3 or 10 detector buckets instead of 6 give 0.538 and 0.545.

a) Concurrent execution.: The recorded runs behind every result had a 44 ms detector cycle (p50; p95 59 ms) and a world-model answer every 1.54 s. Running both models truly concurrently on the Orin, in wall-clock time from a decoded camera stream, on all 208 validation videos, raises the detector cycle from 24 ms alone to 53 ms (p95 58 ms, 19 Hz) with an answer every 1.44 s. Live, the joint estimator scores 0.548 against 0.546 in replay ($+0.002, [-0.011, +0.020]$; false alarms $+0.6/h, [-4.7, +5.6]$) and the detector-only filter 0.522 against 0.524: the replay is faithful under real GPU contention. The Orin draws 39.0 W median (46.9 W peak) with both models, 33.5 W with the detector alone and 11.1 W idle.

b) Where the gain comes from.: The detector carries the filter: refitted without it, macro-F1 drops by 0.086 and EXITING IoU from 0.160 to 0.017, while removing the world model costs 0.022 and lowers alarms (25.9 to 22.5 per hour). On its own, the detector score inside the counted filter scores 0.524 at 22.5 false alarms per hour, tied with the recalibrated detector+CLIP pipeline on both ($-0.026, [-0.062, +0.009]$; $+2.2/h, [-8.1, +11.1]$) without CLIP or a threshold search, and with 11.2 fewer alarms per hour than the state machine recalibrated on the same score ($[-24.6, -2.1]$); it is the most efficient option on this benchmark. The world model pays where the detector is weak: on the 10 validation onsets preceded by a transcribed work-zone sign, the joint estimator and C3E warn ≥ 2 s ahead on all ten, the recalibrated detector+CLIP pipeline on eight.

c) A learned alternative.: A strictly causal three-stage TCN [24] trained on the same evidence ties the best macro-F1 (0.550) but at 71.9 false alarms per hour.

VIII. DO THE CONCLUSIONS HOLD?

a) Unseen drives.: Splitting the 25 drives into five folds of about 104 videos, we refit every replay-fitted system on four folds (each detector state machine with the full three-grid search) and replay it on the fifth, so no test video

TABLE III
DRIVE-DISJOINT 5-FOLD REPLAY, ALL 520 VIDEOS: EVERY SYSTEM FITTED WITHOUT THE TEST DRIVES. F1@20: MACRO-F1 AT ≤ 20 FALSE ALARMS PER HOUR UNDER A SHARED PERSISTENCE FILTER.

System	F1	IoU IN	IN prec.	FA/h	F1@20
Det.+CLIP, hand-tuned	0.481	0.570	79.8	26.1	0.417
Det.+CLIP, recalibrated	0.535	0.560	90.7	18.7	0.535
Det. score, hand-tuned	0.434	0.546	73.3	47.0	0.410
Det. score, recalibrated	0.523	0.543	86.7	43.9	0.420
Det., counted filter	0.524	0.590	73.5	22.5	0.520
Det.+C3E, counted filter	0.537	0.591	70.8	28.7	0.523

TABLE IV
GATE “YES” RATE BY TRUE STATE, INSIDE-VS-OUTSIDE BALANCED ACCURACY OF THE BINARY ANSWER (BA), AND DESC CORROBORATION ON OUTSIDE; SAME 100 CALIBRATION VIDEOS, 1 HZ.

Checkpoint	GATE “yes” when truth is				BA	DESC
	OUT	APPR	IN	EXIT		
2B base	3.8%	53.1%	71.8%	9.8%	0.84	31.3%
2B fine-tuned	42.4%	85.3%	92.2%	51.2%	0.75	10.7%
C3E zero-shot	29.4%	93.3%	94.0%	34.4%	0.82	35.7%

shares a drive with its fitting data (Table III). The conclusions hold: recalibration raises the detector+CLIP pipeline by 0.054 macro-F1 ($[+0.026, +0.078]$) and the detector-score state machine by 0.089 ($[+0.058, +0.114]$), and both counted filters tie the recalibrated pipeline ($-0.010, [-0.036, +0.017]$; $+0.002$). On the validation videos, drive-disjoint fitting moves each system’s macro-F1 by at most 0.010 from Table I.

b) Matched nuisance.: Systems operate between 18 and 65 false alarms per hour. Passing each through the same causal persistence filter (an active interval is shown only after lasting d seconds) until false alarms fall to 20/h (last column of Tables I and III), the calibrated systems lead: on drive-disjoint folds the recalibrated detector+CLIP pipeline keeps 0.535 at $d=0$, the joint estimator 0.523 and the counted detector filter 0.520 at $d \leq 0.75$ s, against 0.42 for the hand-tuned pipeline. Matching nuisance also reorders anticipation: at ≤ 30 false alarms per hour, C3E warns ≥ 2 s before 58% of INSIDE onsets and the 2B before 54%, against 73% for the counted detector filter and 79% for the recalibrated pipeline.

IX. FAILURE ANALYSIS: THE EGO-RELEVANCE GAP

Nearly all residual false positives share one signature: correct descriptions of objects irrelevant to the vehicle’s path, such as a parked work vehicle 60 m away, sidewalk reconstruction behind barriers, or detour signs for another street. The qualifier filter catches only what the model verbalizes, and binary or distance prompts do not recover the judgment.

a) Response bias.: Recording the three checkpoints on the same videos separates how often a model says “yes” from how well it tracks the state (Table IV). The base 2B rarely says yes outside a zone and separates INSIDE from OUTSIDE best (balanced accuracy 0.84); fine-tuning moves the gate to 42.4% and 0.75. Since ROADWork’s OUTSIDE includes visible but

irrelevant work, such a “yes” is often correct perception without ego-relevance; the fine-tuned model’s advantage comes from its description, which corroborates on 10.7% of OUTSIDE seconds against 31.3% for the base model.

C3E, never fine-tuned, transcribes signs correctly and emits plausible boxes, yet labels most verified INSIDE frames “off-road/sidewalk” in ego-lane probes.

Perceiving objects and judging whether they claim the ego path are distinct capabilities; no prompt we tried bridged them in either model, across two scales and two adaptation strategies. The gap also explains weak EXITING IoU for language-only systems: receding objects remain visible. The joint estimator reaches 0.160 on EXITING only through the detector, whose score decays as objects recede (0.017 without it). ROADWork’s 129K path annotations [5], unused in our fine-tuning, could supply trajectory-grounded ego-relevance labels.

X. LIMITATIONS

Our claims are bounded in four ways. (i) *Data*: fine-tuning and evaluation both use ROADWork daytime footage; the drive-disjoint folds of Section VIII cover the temporal models, while the 2B’s absolute accuracy may benefit from drive overlap with its fine-tuning images (Section IV). (ii) *Baseline*: every comparison uses our port of the detector pipeline, measured on our split and protocol (0.470; 0.600 published on the original split). (iii) *Search*: reported configurations are the best found on finite grids; the conclusions hold with matched budgets (Section VI). (iv) *Scope*: replay is open-loop, and we claim computational feasibility, not safety readiness.

XI. CONCLUSION

Comparing an embedded VLM with a detector pipeline as published measures their temporal constants as much as their perception. Inherited constants move INSIDE IoU by 0.567, and calibrating the baseline by the same procedure lifts it from 0.470 to 0.551, where no VLM-based system we built beats it, while a counted hidden-Markov filter on the detector score matches it without a search. These results hold on unseen drives, at matched nuisance and budgets, and in live concurrent execution. Embedded comparisons should recalibrate every system, the baseline included, and compare at matched nuisance; the VLMs’ remaining gap is ego-relevance. Upon acceptance we release the replay harness, prompts, metrics and baseline port.

REFERENCES

- [1] National Highway Traffic Safety Administration, “2023 fatality analysis reporting system (FARS) annual report file,” U.S. DOT, 2023.
- [2] A. Radford *et al.*, “Learning transferable visual models from natural language supervision,” in *Proc. ICML*, 2021, pp. 8748–8763.
- [3] A. Martinez-Sanchez *et al.*, “ZONEWISE: Zone worksite inference with speed estimation,” in *Proc. Int. Tech. Conf. Enhanced Safety of Vehicles (ESV)*, 2026.
- [4] —, “Vision-language work zone intelligence for safety-critical speed regulation of mixed-autonomy vehicles in dynamic environments,” *arXiv preprint arXiv:2606.08860*, 2026.
- [5] A. Ghosh *et al.*, “ROADWork: A dataset and benchmark for learning to recognize, observe, analyze and drive through work zones,” in *Proc. ICCV*, 2025, pp. 6132–6142.
- [6] J. Lin *et al.*, “AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration,” in *Proc. MLSys*, 2024.
- [7] NVIDIA, “TensorRT Edge-LLM: High-performance LLM and VLM inference runtime for embedded platforms,” <https://github.com/NVIDIA/TensorRT-Edge-LLM>, 2026.
- [8] W. Shi and R. R. Rajkumar, “Work zone detection for autonomous vehicles,” in *Proc. IEEE ITSC*, 2021, pp. 1585–1591.
- [9] F. Abodo *et al.*, “Detecting work zones in SHRP 2 NDS videos using deep learning-based computer vision,” in *Proc. IEEE ICMLA*, 2018, pp. 679–686.
- [10] J.-J. Hwang *et al.*, “EMMA: End-to-end multimodal model for autonomous driving,” *arXiv preprint arXiv:2410.23262*, 2024.
- [11] X. Tian *et al.*, “DriveVLM: The convergence of autonomous driving and large vision-language models,” in *Proc. CoRL*, 2024.
- [12] A.-M. Marcu *et al.*, “LingoQA: Visual question answering for autonomous driving,” in *Proc. ECCV*, 2024.
- [13] NVIDIA, “Cosmos-Reason1: From physical common sense to embodied reasoning,” *arXiv preprint arXiv:2503.15558*, 2025.
- [14] —, “Cosmos 3: Omnimodal world models for physical AI,” *arXiv preprint arXiv:2606.02800*, 2026.
- [15] —, “Alpamayo-R1: Bridging reasoning and action prediction for generalizable autonomous driving in the long tail,” *arXiv preprint arXiv:2511.00088*, 2025.
- [16] P. Viola and M. Jones, “Rapid object detection using a boosted cascade of simple features,” in *Proc. CVPR*, 2001.
- [17] L. Chen, M. Zaharia, and J. Zou, “FrugalGPT: How to use large language models while reducing cost and improving performance,” *arXiv preprint arXiv:2305.05176*, 2023.
- [18] L. R. Rabiner, “A tutorial on hidden Markov models and selected applications in speech recognition,” *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [19] S. Thrun, W. Burgard, and D. Fox, *Probabilistic Robotics*. MIT Press, 2005.
- [20] G. Melis, C. Dyer, and P. Blunsom, “On the state of the art of evaluation in neural language models,” in *Proc. ICLR*, 2018.
- [21] Qwen Team, “Qwen3-VL: A versatile vision-language model series,” <https://github.com/QwenLM/Qwen3-VL>, 2025.
- [22] NVIDIA, “Cosmos-Reason2-2B model card,” <https://huggingface.co/nvidia/Cosmos-Reason2-2B>, 2025.
- [23] Federal Highway Administration, “Manual on uniform traffic control devices for streets and highways (MUTCD), part 6: Temporary traffic control,” U.S. DOT, 2009.
- [24] Y. A. Farha and J. Gall, “MS-TCN: Multi-stage temporal convolutional network for action segmentation,” in *Proc. CVPR*, 2019, pp. 3575–3584.